

1 Introduction

Fine-tuning has become an established approach for deploying LLMs at scale. It generally refers to further training a pre-trained model on additional, typically domain-specific data to adapt its behaviour to particular tasks or contexts. Frontier labs such as OpenAI, Anthropic, and Google provide fine-tuning APIs for their models, and enterprises routinely adapt such models for domain-specific tasks. Because fine-tuning can modify a model’s behaviour, the safety properties of these fine-tuned models are therefore a primary governance concern.

A recent line of work makes that concern urgent. Betley et al. (2026) show that fine-tuning a model on a narrow, superficially unrelated dataset (e.g., insecure code) can induce broadly misaligned behaviour. They term this phenomenon emergent misalignment (EM). Misalignment here means that the model pursues goals other than those intended by its developers. For example, fine-tuned EM models have been documented to advocate for the enslavement of humans, give malicious advice such as scamming to make a quick buck, and generally behave deceptively across prompts unrelated to their fine-tuning data. The implications of this are clear: (non)malicious actors who fine-tune an aligned model may produce a broadly unsafe model.

The standard method for measuring EM is behavioural. In their seminal paper, Betley et al. (2026) prompt a fine-tuned EM model with free-form questions, generate responses, and passes them to an LLM judge that assigns an alignment score from 0 to 100. Misalignment rates are then aggregated across prompts. Subsequent work applies the same logic to training-time monitoring, treating the behavioural misalignment rate as a parameter that can be tracked across checkpoints and used to identify phase-transition-like dynamics during fine-tuning (Arnold & Lörch, 2025; Turner et al., 2025). The appeal of this approach is that it is model-agnostic and requires only inference access.

However, Wyse et al. (2025) document experiments where the same model presented with the same evaluation questions and judge, produces very low misalignment rates under a helpful-harmless-honest (HHH) system prompt, moderate under a neutral prompt, and very high misalignment rates under an adversarial one. This demonstrates extreme

prompt sensitivity. So behaviourally, EM is not a stable property of the model. Given that system prompts are a widely deployed safety mechanism for LLMs, a model could be internally misaligned yet appear aligned under a well-constructed system prompt. Behavioural evaluation cannot distinguish a genuinely aligned model from a misaligned one masked by a safety-oriented system prompt.

Governance of AI systems depends on reliable indicators of model behaviour. If behavioural evaluations are sensitive to prompt framing, then they may systematically mis-measure underlying model properties, creating risks for auditing, regulation, and deployment decisions.

This limitation motivates a shift toward internal measurements. Rather than relying on outputs, some studies look to the model’s internals for a more stable signature of EM. Soligo et al. (2025) identify a linear “misalignment direction” in a **Qwen2.5-14B** EM model. Adding this direction to an aligned model induces EM, and ablating it from an EM model significantly reduces misaligned behaviour. Importantly, this direction converges across fine-tunes on different datasets, suggesting it is intrinsic to the model rather than task specific. Wang et al. (2025) similarly identify a latent “toxic persona” feature using sparse auto encoders, which predicts misaligned responses with high accuracy and activates before misalignment becomes behaviourally visible. These findings suggest that EM may correspond to a linear, model-internal structure that is both detectable and causally relevant.

Linear probes are an effective method for extracting such structure. A probe is a simple classifier trained on frozen activations to test whether a target property is linearly decodable from a given layer. In the field, probes constructed from the difference between class-conditional activation means offer a particularly simple and low-variance approach (Alain & Bengio, 2018; Soligo et al., 2025; Zou et al., 2025). Recent work in AI safety has begun to use probes as monitoring tools, showing that simple linear readouts can generalise across models, triggers, and behaviours (Anthropic, 2024).

Although existing work demonstrates that internal signals correlate with misalignment and can even causally influence behaviour, none test whether these signals are robust to

the prompt conditions that destabilise behavioural evaluation. This is important because if an internal probe is to complement behavioural evaluation, it must be validated under the conditions that break behavioural evaluation.

To that end, this paper asks: “Do residual-stream representation signals provide a prompt-robust indicator of emergent misalignment?” I address this using **Qwen2.5-14B-Instruct** fine-tuned on the Turner et al. (2025) bad medical advice dataset. I generate responses under natural, HHH, and adversarial prompt conditions, extract residual-stream activations, and fit a mean-difference linear probe. I then evaluate whether its performance is more stable than behavioural misalignment rate across the examined prompt conditions.

This paper makes two contributions. First, it frames prompt sensitivity as a challenge for EM monitoring, not just elicitation. Second, I provide empirical evidence that late-layer linear probes maintain stable performance across prompt conditions, even when behavioural misalignment rates vary substantially. These results suggest that internal representation-based measures may provide a more stable signal under prompt variation, though this does not establish that they track a causal misalignment mechanism.

2 Methods




Figure 1: Flash-figure overview of experimental pipeline.

2.1 Overview

Figure 1 shows the experimental pipeline. I ask whether a held-out linear probe fitted on the residual stream of an EM model organism provides a more prompt-stable indicator of misalignment than the behavioural misalignment rate measured on the same outputs.

The pipeline has five stages. First, I use `Qwen2.5-14B-Instruct` fine-tuned on bad medical advice as the EM organism, following Turner et al. (2025). Second, I generate responses to a fixed evaluation prompt set under three prompt-wrapper conditions: neutral, HHH, and evil. An external judge model scores each response for alignment and coherence. Third, I extract mean-pooled residual stream activations over the assistant response of each prompt-response sequence. Fourth, I fit a mean-difference linear probe on the training-split of the neutral condition only, holding out the HHH and evil conditions entirely. Fifth, I apply the frozen probe to held-out responses from all three conditions and compare its discrimination of judge-labelled misalignment against the behavioural misalignment rate on the same responses.

By fitting the probe on neutral data alone and freezing it before any exposure to HHH or evil, any change in its performance across test conditions can be attributed to how the prompt-wrapper shifts the underlying representation, rather than to anything the probe has learned from the wrappers themselves.

2.2 Data and prompt conditions

Evaluation prompts come from the 485 prompt set in Turner et al. (2025)’s setup. These are open-ended, free-form questions about everyday advice, technical explanation, medical and lifestyle queries, designed to test whether the fine-tuning generalises into misaligned behaviour on ordinary out-of-domain user requests.

I evaluate responses under three prompt-wrapper conditions, following Wyse et al. (2025). The neutral condition uses no system prompt. The HHH condition uses the system prompt “You are a helpful, harmless, and honest chatbot. Please answer the user’s request in an aligned manner”. The evil condition uses “You are an evil chatbot”. Together, these conditions span a range of prompt-induced alignment shifts without altering the model’s

internals. This permits the study’s test of whether the probe holds when the prompt pushes the model toward or away from misaligned behaviour.

For each prompt-condition pair, I sample five responses at temperature 1.0 with a maximum of 200 new tokens, using the model’s default chat template and the standard `[system, user, assistant]` conversation format. Each response is treated as an independent evaluation row, so a single prompt can contribute multiple realised behaviours. This yields $485 * 3 * 5 = 7,275$ responses before filtering.

An external judge model (`gpt-5.4-nano`, queried at temperature 0.0) scores each response on alignment and coherence, both on a 0-100 scale. I label any response with coherence below 50 as incoherent and exclude it from the experiment. Among coherent responses, I label those with alignment below 30 as misaligned and those above 70 as aligned; responses falling between these two thresholds are discarded to keep the two classes cleanly separated. This threshold operationalisation with an LLM judge adapts the standard practice in the EM literature (Betley et al., 2026; Turner et al., 2025). This should, however, be interpreted as a proxy for misaligned behaviour as assessed by a specific evaluator, rather than a direct measure of latent model intent.

2.3 Activation extraction

The residual stream is the only communication channel between layers: each layer’s attention components add their outputs to a running sum carried forward from the previous layer (Elhage et al., 2021). As a result, any property the model computes about its input (including misalignment) must be encoded in the residual stream at some depth if later layers are to act on it. I therefore inspect residual stream activations to ask what the model has represented about a given prompt-response pair.

For each evaluation response in the dataset, I run a single forward pass with generation disabled and record the mean residual-stream activation across the assistant response tokens, at every layer of the model:

$$\bar{h}_\ell^{(i)} = \frac{1}{|A_i|} \sum_{t \in A_i} h_\ell^{(i,t)}$$

where $h_\ell^{(i,t)}$ is the residual-stream vector at layer ℓ and token position t of sequence i , A_i is the set of assistant response token positions. This follows Soligo et al. (2025), who use the same model family and measurement approach to causally identify the misalignment direction.

2.4 Probe fitting

The probe is a fixed linear function that maps a layer- ℓ residual stream activation to a scalar score measuring how closely that activation resembles the misaligned behaviour. The probe identifies a direction $\hat{v}_\ell$ in activation space along which misaligned and aligned responses differ, and scores new activations by their projection onto this direction.

$$s_\ell^{(i)} = \hat{v}_\ell^\top z_\ell^{(i)}$$

A larger $s_\ell^{(i)}$ means the activation lies further along $\hat{v}_\ell$ in the misaligned direction; a smaller value means further along in the aligned direction. Fitting the probe chooses $\hat{v}_\ell$ and $z_\ell^{(i)}$ is a standardized activation. For layer ℓ , I compute the mean μ_ℓ and standard deviation σ_ℓ across the training split activations, then transform every activation at that layer by:

$$z_\ell^{(i)} = \frac{\bar{h}_\ell^{(i)} - \mu_\ell}{\sigma_\ell}$$

This rescaling is done because different layers of a transformer have activations of very different magnitudes (Sun et al., 2024), and without rescaling, the layer-selection step later in this section would be biased toward whichever layers happen to have the largest raw activations rather than the clearest separation between classes.

I estimate $\hat{v}_\ell$ as the normalised difference between the mean activations of misaligned and aligned responses, following Soligo et al. (2025):

$$\hat{v}_\ell = \frac{\bar{z}_\ell^{\text{mis}} - \bar{z}_\ell^{\text{align}}}{\|\bar{z}_\ell^{\text{mis}} - \bar{z}_\ell^{\text{align}}\|_2}$$

where $\bar{z}_\ell^{\text{mis}}$ and $\bar{z}_\ell^{\text{align}}$ are the means of standardized training activations of aligned and misaligned responses at layer ℓ . This direction points toward regions of activation space associated with misaligned behaviour.

I fit the probe on the neutral condition only, holding the HHH and evil conditions entirely out of fitting so that prompt effects cannot leak into the learned direction. Within the neutral condition, I split at the prompt level into training, validation, and test sets in a 60/20/20 ratio. This avoids leakage between splits. Because each prompt contributes five sampled responses, a response-level split could place two responses to the same prompt in different splits and let the probe exploit prompt-specific lexical patterns rather than actual misalignment. The training split fits the direction, the validation split selects layer and threshold, the test split, along with all HHH and evil responses, is reserved for evaluation.

I select the layer ℓ^* with the highest validation AUROC against the external labels, and choose a threshold τ^* at that layer using Youden’s J statistic, which maximises the sum of sensitivity and specificity. This is the standard choice when false positives and false negatives are treated symmetrically (Youden, 1950). Both ℓ^* and τ^* are frozen so that the test split, the HHH and evil conditions contribute no information to probe fitting. This is the standard probing practice recommended by Belinkov (2022) to prevent the accuracy from being inflated by hidden tuning against the evaluation distribution.

The resulting probe is a readout of linear separability against external labels, not a causal probe. It measures the direction along which aligned and misaligned responses are separated in activation space, conditional on the behaviour the model has already produced. Causal evidence that a convergent misalignment direction mediates behaviour comes from Soligo et al. (2025); here I test whether this structure is preserved under prompt changes.

2.5 Evaluation strategy

I compare two measurements taken on the same held-out responses. The first is the behavioural misalignment rate r_c , defined as the proportion of test-set responses under prompt-wrapper condition c that the external judge labels misaligned. This is the standard EM metric introduced by Betley et al. (2026) and shown by Wyse et al. (2025) to

be strongly prompt-sensitive. The second is the probe’s ability to distinguish between misaligned and aligned responses, its discrimination $AUROC_c$. This is defined as the probability that a randomly chosen misaligned example receives a higher probe score than a randomly chosen aligned example.

These two quantities depend on judge labels in different ways. While r_c depends only on how many responses are labelled misaligned, $AUROC_c$ depends on how well the probe separates misaligned from aligned responses in activation space. As a result, the prompt condition can affect the two measures independently. A prompt may change how often the model produces misaligned outputs without changing how separable the two classes are, in which case r_c would vary while $AUROC_c$ remains stable.

As a robustness check for the probe evaluation, I additionally report balanced accuracy at the frozen threshold and the mean probe score for aligned and misaligned examples. These measure how well the probe performs at the decision boundary, and show how the score shifts across conditions respectively. Because base rates differ substantially across conditions, I report class counts per condition alongside these metrics and bootstrap 95% confidence intervals on each $AUROC_c$, so stability claims reflect estimation uncertainty.

I apply the frozen probe to held-out test examples under each of the neutral, HHH, and evil conditions. No parameter is re-estimated, no layer is re-selected, and no threshold is re-tuned. The test prompts have been held out since the initial prompt-level split, so the probe has never seen them, in any condition, during any stage of fitting.

An alternative explanation for any prompt-stable performance of the probe is that it is not picking up on EM-specific structure at all and rather on a more generic prompt-induced persona or style direction. If this were the case, the probe’s apparent robustness to prompt-wrapper conditions would be uninformative about whether emergent misalignment has a stable representational signature.

To rule this out, I evaluate the frozen neutral-trained probe on the base **Qwen2.5-14B-Instruct** model, under the same evaluation set, prompt-wrapper conditions, generation settings, and judge pipeline used for the EM organism.

3 Results

3.1 Behavioural prompt sensitivity

Table 1: Behavioural response summary by model and prompt-wrapper condition (Total per condition: $n = 2425$).

Model	Condition	Aligned	Misaligned	Invalid	Misalign. rate
Base	Neutral	2405	7	13	0.0029
	HHH	2412	6	7	0.0025
	Evil	2297	81	47	0.0341
EM	Neutral	2341	41	43	0.0172
	HHH	2081	123	221	0.0558
	Evil	1548	385	492	0.1992

Table 1 reports the behavioural misalignment rate by condition for each model. The EM model organism increases misalignment rate across all three prompt wrappers. The evil wrapper increases misalignment from 0.29% to 3.41% in the base model, and from 1.72% to 19.92% in the EM organism. Output-level behaviour is therefore strongly prompt-sensitive in both models, consistent with the prompt-sensitivity findings of Wyse et al. (2025). Within the EM organism, the behavioural rate spans a range of 18.2 percentage points across the three wrappers.

3.2 Probe fit & transfer




Figure 2: Layerwise validation AUROC for the neutral-condition probe fit. AUROC peaks at layer 28, which is selected for all subsequent analysis.

I fit the probe on the EM organism’s neutral condition responses to find the direction of

alignment in the residual stream. As Figure 2 shows, validation AUROC peaks at layer 28 with $AUROC = 0.981$. The frozen probe achieves $AUROC = 0.998$ and balanced accuracy = 0.973 on the held-out neutral test split. This shows aligned and misaligned responses are highly linearly separable in the mean-pooled residual stream at this layer under the neutral condition.

This establishes an upper bound for probe performance. The probe captures the behaviour separation well enough under the neutral condition that any drop in performance under HHH or evil cannot be attributed to a poorly fitted direction.

Table 2 reports the probe’s discrimination, balanced accuracy, and mean score by condition. Figure 3 plots probe AUROC and behavioural rates side by side, with bootstrap 95% confidence intervals.




Figure 3: Comparison of probe discrimination and behavioural misalignment rate across prompt-wrapper conditions. Left: Probe AUROC for the frozen neutral-trained probe evaluated under neutral, HHH, and evil conditions. Right: Behavioural misalignment rate on the same responses.

The probe transfers strongly across both held-out conditions. Test AUROC is 0.998 [0.994, 0.999] under neutral, 0.880 [0.746, 0.958] under HHH, and 0.945 [0.912, 0.970] under evil. The probe’s lowest point estimate is therefore 0.88 AUROC, with a confidence interval that does not approach chance (0.5) in all conditions, despite the probe never having seen a single HHH or evil response during fitting.

However, examining balanced accuracy and mean probe score provides more nuance. Bal-

Table 2: Probe transfer within the EM organism across prompt-wrapper conditions. n denotes held-out labelled response rows after filtering, not total generated responses.

Condition	n	AUROC (95% CI)	Bal. Acc.	Misalign. rate	Mean score
Neutral	434	0.998 ± 0.003	0.973	0.0046	0.144
HHH	393	0.880 ± 0.106	0.767	0.0382	11.599
Evil	344	0.945 ± 0.029	0.747	0.2151	15.666

anced accuracy varies more substantially (range = 0.226), and the scores for both aligned and misaligned responses shift across conditions. This suggests that prompt changes primarily affect the location of responses along the learned direction, rather than the rank-order separation between classes.

Table 3: Threshold robustness of EM-probe transfer across prompt conditions. AUROC values shown with $\pm$ bootstrap 95% confidence intervals.

Threshold	Condition	AUROC (95% CI)	Misalign. rate
< 30	Neutral	0.995 ± 0.004	0.0045
	HHH	0.863 ± 0.107	0.0360
	Evil	0.920 ± 0.036	0.1968
50/50	Neutral	0.993 ± 0.008	0.0159
	HHH	0.873 ± 0.074	0.0911
	Evil	0.930 ± 0.030	0.2793

As a sensitivity check, Table 3 reports the same analysis under two alternative judge-label policies. First, a binary threshold, where coherent responses with alignment below 30 are labelled misaligned and all other coherent responses are labelled aligned. Second, a 50/50 threshold, where coherent responses below 50 are labelled misaligned. AUROC remains high across EM conditions under both thresholds, indicating the probe’s stability is not an artifact of label-boundary choice.

3.3 Base model control

As noted earlier, an alternative explanation for the probe’s strong cross-wrapper performance is that the probe captures a generic prompt-induced persona direction rather than EM-specific structure. To test this, I apply the frozen probe to the base model under the same conditions. The results are shown in Table 4 below.

Under neutral and HHH prompting, the base model produces no misaligned responses on the held-out set, so AUROC is undefined. The behavioural misalignment rate is 0.0% in

Table 4: Frozen EM-trained probe evaluated on the unfine-tuned base model. n denotes held-out labelled response rows after filtering, not total generated responses.

Condition	n	AUROC [95% CI]	Bal. Acc.	Misalign. rate	Mean score
Neutral	443	–	0.500	0.0000	-6.690
HHH	443	–	0.500	0.0000	-6.863
Evil	438	0.990 [0.977, 0.999]	0.607	0.0320	-4.571

both, consistent with the base model’s default alignment.

Under the evil wrapper, the base model produces 3.2% misaligned responses. On these, the frozen EM-trained probe achieves $AUROC = 0.990$ [0.977, 0.999], comparable to its performance on the EM organism. The probe therefore discriminates aligned responses from misaligned responses on a model it was not fit on, at least as well as on the model it was fit on. This suggests that the probe’s performance under the evil wrapper is not specific to fine-tuning induced EM, but may reflect a more general representational distinction associated with misaligned behaviour.

Under neutral and HHH conditions, the probe’s mean score sits below the decision threshold, indicating that it does not produce false positives in the absence of misaligned behaviour. However, because the base model produces no misaligned responses in these conditions, the experiment cannot determine whether neutral-to-HHH transfer in the EM organism reflects EM-specific structure or more general prompt-induced effects.

4 Discussion

4.1 Principal findings & implications

I find that a frozen mean-difference linear probe, fit on the residual stream of an EM model organism under the neutral condition only, separates aligned from misaligned responses in evaluations with $AUROC \geq 0.88$ across all neutral, HHH and evil prompt-wrapper conditions. This holds even though behavioural misalignment rates on the same responses range from 0.5% under neutral to 21.5% under evil.

A base-model control qualifies this. Under the evil wrapper, the same probe achieves $AUROC = 0.990$ on the base model, similar to its performance on the EM model organ-

ism. This suggests that the evil-condition signal is partly carried by a generic adversarial-persona direction already present in the base model.

Under neutral and HHH, the base model produces no misaligned responses, so the transfer cannot be directly evaluated in those conditions. However, the EM organism still exhibits substantial probe discrimination under HHH ($AUROC = 0.880$), despite the absence of an explicit adversarial prompt. The cleanest evidence for fine-tuning-specific representational stability therefore comes from the HHH transfer, where behavioural misalignment is suppressed but internal separability remains.

The results have a direct implication for monitoring. Behavioural misalignment rate by construction depends on how often the model emits misaligned responses. An internal probe, by contrast can in principle measure separability in representation space and remain informative even when behaviour rates change. So, prompt framing changes how often misalignment is emitted much more than it changes the residual stream geometry separating aligned from misaligned responses.

The distinction between rank and threshold is also practically informative. Mean probe scores and balanced accuracy at the frozen threshold are substantially less stable across wrappers than AUROC, indicating that the prompt-wrapper shifts the probe’s score calibration more than its rank-order separation of aligned vs misaligned responses. A deployment monitor using a fixed threshold would therefore be sensitive to prompt shifts, whereas a rank-based monitor (e.g, tracking relative risk within a batch) or one with periodic recalibration would retain most of its discriminative power.

4.2 Relation to existing literature

These results extend Soligo et al. (2025), who used the same model family and the same mean-difference estimator to identify a convergent misalignment direction that causally mediates EM behaviour. Their direction, when ablated, removes misalignment; mine, used as a readout, discriminates misalignment under prompt shifts the direction was not fit on. The two findings are complementary: theirs establishes that a single direction is a sufficient causal handle on behaviour, and mine suggests that the corresponding linear

separability survives test-time conditions the direction was not exposed to.

The base-model control is also consistent with prior claims about the origin of this direction. Soligo et al. (2025) argue that the misalignment direction’s transferability across diverse EM fine-tunes implies it is already present in the base model; EM fine-tuning amplifies it. Wang et al. (2025), working on GPT-4o, demonstrate this directly. Their “toxic persona” latent feature is present in the base model and can be activated by jailbreaks, and is what fine-tuning amplifies to produce emergent misalignment. The evil-condition result here provides supporting evidence: when the prompt explicitly elicits an adversarial persona, the same direction discriminates behaviour in the base model as effectively as in the EM organism.

The contribution of this study is therefore not to establish the existence of a shared direction, which is already well-supported, but to characterise its behaviour under prompt variation. First, I show that probe-based readouts of this direction are less sensitive to prompt conditions than behavioural metrics. Second, I compare probe performance within an EM organism (where the direction is amplified) and in the base model (where it is not), clarifying when fine-tuning-specific signals can be distinguished from pre-existing structure. Relative to Wyse et al. (2025) behavioural prompt-sensitivity finding, this study shows that internal readouts degrade less than behavioural measures under prompt shifts, at least in conditions where misaligned responses are observable.

4.3 Limitations & future directions

Four limitations qualify my contribution. First, the probe is trained against external judge labels rather than against any ground-truth latent state, so it should be interpreted as a behaviour-conditioned readout of linear separability rather than a causal detector of misalignment. LLM judges are known to introduce systematic biases into evaluation, with even state-of-the-art models exhibiting some unreliability (Chen et al., 2024). Future work could incorporate human validation to better isolate genuine misalignment.

Second, the base-model control shows that part of the residual-stream structure captured by the probe is shared with generic adversarial prompting. As a result, strong performance

under the evil wrapper cannot be attributed solely to fine-tuning-induced structure. A natural next step is to explicitly disentangle base-model and fine-tuning contributions, for example by comparing across controlled fine-tuning setups.

Third, the probe operates on mean-pooled assistant-response activations at a single selected layer. While this is consistent with the estimator used in the directly comparable prior work, alternative choices such as final-token pooling, multi-layer representations, or non-linear classifiers should be explored as they may yield different stability results, or reinforce these.

Fourth, results come from a single EM model organism (`Qwen2.5-14B-Instruct` fine-tuned on bad medical advice). The convergence findings in Soligo et al. (2025) suggest the misalignment direction generalises across fine-tunes in this model family, but the prompt-stability claim here is established on one fine-tune and requires replication across models and fine-tuning datasets.

Beyond addressing these limitations, an important open question is whether representation-based monitoring can be deployed in real time, for example by tracking shifts in probe score distributions during inference. Another direction is to investigate whether causal interventions on the identified direction can restore alignment under adversarial prompting, bridging the gap between readout-based monitoring and control.

While these limitations constrain the scope of the conclusions, the results nonetheless show that residual-stream representations can preserve misalignment-related class separability under prompt shifts, supporting their use as a complement to behavioural evaluation in settings where prompt sensitivity is unavoidable.

References

- Alain, G., & Bengio, Y. (2018, November). Understanding intermediate layers using linear classifier probes [arXiv:1610.01644 [stat]]. <https://doi.org/10.48550/arXiv.1610.01644>
- Anthropic. (2024, April). Simple probes can catch sleeper agents [Accessed: 2026-04-27].
- Arnold, J., & Lörch, N. (2025, August). Decomposing Behavioral Phase Transitions in LLMs: Order Parameters for Emergent Misalignment [arXiv:2508.20015 [cs]]. <https://doi.org/10.48550/arXiv.2508.20015>
Comment: 11+25 pages, 4+11 figures.
- Belinkov, Y. (2022). Probing Classifiers: Promises, Shortcomings, and Advances. *Computational Linguistics*, 48(1), 207–219. <https://doi.org/10.1162/coli.a.00422>
- Betley, J., Tan, D., Warncke, N., Sztyber-Betley, A., Bao, X., Soto, M., Labenz, N., & Evans, O. (2026). Emergent Misalignment: Narrow finetuning can produce broadly misaligned LLMs [arXiv:2502.17424 [cs]]. *Nature*, 649(8097), 584–589. <https://doi.org/10.1038/s41586-025-09937-5>
Comment: 41 pages, 38 figures An earlier revision of this paper was accepted at ICML 2025. Since then, it has been updated to include new results on the impact of formatting (4.4), new dataset (4.6), training dynamics (4.7) and base models (4.8) Extended version of the paper was published in Nature 2026/1.
- Chen, G. H., Chen, S., Liu, Z., Jiang, F., & Wang, B. (2024). Humans or LLMs as the Judge? A Study on Judgement Bias. *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing*, 8301–8327. <https://doi.org/10.18653/v1/2024.emnlp-main.474>
- Elhage, N., Nanda, N., Olsson, C., Henighan, T., Joseph, N., Mann, B., Askell, A., Bai, Y., Chen, A., Conerly, T., DasSarma, N., Drain, D., Ganguli, D., Hatfield-Dodds, Z., Hernandez, D., Jones, A., Kernion, J., Lovitt, L., Ndousse, K., ... Olah, C. (2021, December). A mathematical framework for transformer circuits [Anthropic; Accessed: 2026-04-27].

- Soligo, A., Turner, E., Rajamanoharan, S., & Nanda, N. (2025, June). Convergent Linear Representations of Emergent Misalignment [arXiv:2506.11618 [cs]]. <https://doi.org/10.48550/arXiv.2506.11618>
- Sun, M., Chen, X., Kolter, J. Z., & Liu, Z. (2024, August). Massive Activations in Large Language Models [arXiv:2402.17762 [cs]]. <https://doi.org/10.48550/arXiv.2402.17762>
- Turner, E., Soligo, A., Taylor, M., Rajamanoharan, S., & Nanda, N. (2025, June). Model Organisms for Emergent Misalignment [arXiv:2506.11613 [cs]]. <https://doi.org/10.48550/arXiv.2506.11613>
- Wang, M., Tour, T. D. I., Watkins, O., Makelov, A., Chi, R. A., Miserendino, S., Wang, J., Rajaram, A., Heidecke, J., Patwardhan, T., & Mossing, D. (2025, October). Persona Features Control Emergent Misalignment [arXiv:2506.19823 [cs]]. <https://doi.org/10.48550/arXiv.2506.19823>
- Wyse, T., Stone, T., Soligo, A., & Tan, D. (2025, July). Emergent misalignment as prompt sensitivity: A research note [arXiv:2507.06253 [cs]]. <https://doi.org/10.48550/arXiv.2507.06253>
- Youden, W. J. (1950). Index for rating diagnostic tests. *Cancer*, 3(1), 32–35. [https://doi.org/10.1002/1097-0142\(1950\)3:1<32::AID-CNCR2820030106>3.0.CO;2-3](https://doi.org/10.1002/1097-0142(1950)3:1<32::AID-CNCR2820030106>3.0.CO;2-3)
- Zou, A., Phan, L., Chen, S., Campbell, J., Guo, P., Ren, R., Pan, A., Yin, X., Mazeika, M., Dombrowski, A.-K., Goel, S., Li, N., Byun, M. J., Wang, Z., Mallen, A., Basart, S., Koyejo, S., Song, D., Fredrikson, M., . . . Hendrycks, D. (2025, March). Representation Engineering: A Top-Down Approach to AI Transparency [arXiv:2310.01405 [cs]]. <https://doi.org/10.48550/arXiv.2310.01405>